

# Name:Viviyan Richards W

Roll no:205229133

## Lab3. Computing Document Similarity using VSM

### EXERCISE-1: Print TFIDF values

```
In [1]: from sklearn.feature_extraction.text import TfidfVectorizer
```

```
In [2]: import pandas as pd
```

```
In [3]: docs = ["good movie", "not a good movie", "did not like", "i like it", "good one"]
```

```
In [4]: tfidf = TfidfVectorizer(min_df=2, max_df=0.5, ngram_range=(1, 2))
features = tfidf.fit_transform(docs)
print(features)
```

```
(0, 0)      0.7071067811865476
(0, 2)      0.7071067811865476
(1, 3)      0.5773502691896257
(1, 0)      0.5773502691896257
(1, 2)      0.5773502691896257
(2, 1)      0.7071067811865476
(2, 3)      0.7071067811865476
(3, 1)      1.0
```

```
In [5]: df = pd.DataFrame(
    features.todense(),
    columns=tfidf.get_feature_names())
print(df)
```

	good movie	like	movie	not
0	0.707107	0.000000	0.707107	0.000000
1	0.577350	0.000000	0.577350	0.577350
2	0.000000	0.707107	0.000000	0.707107
3	0.000000	1.000000	0.000000	0.000000
4	0.000000	0.000000	0.000000	0.000000

### EXERCISE-2:

1. Change the values of min\_df and ngram\_range and observe various outputs

```
In [6]: tfidf = TfidfVectorizer(min_df=1, max_df=0.6, ngram_range=(1, 2))
features = tfidf.fit_transform(docs)
print(features)
```

```
(0, 3)      0.6098184563533858
(0, 8)      0.6098184563533858
(0, 2)      0.5062044059286201
(1, 10)     0.5422255279709232
(1, 9)      0.4374641418373903
(1, 3)      0.4374641418373903
(1, 8)      0.4374641418373903
(1, 2)      0.36313475547801904
(2, 11)     0.4821401170833009
(2, 1)      0.4821401170833009
(2, 6)      0.3889876106617681
(2, 0)      0.4821401170833009
(2, 9)      0.3889876106617681
(3, 7)      0.6141889663426562
(3, 5)      0.6141889663426562
(3, 6)      0.49552379079705033
(4, 4)      0.6390704413963749
(4, 12)     0.6390704413963749
(4, 2)      0.42799292268317357
```

```
In [7]: df = pd.DataFrame(
    features.todense(),
    columns=tfidf.get_feature_names())
print(df)
```

	did	did not	good	good movie	good one	it	like \
0	0.00000	0.00000	0.506204	0.609818	0.00000	0.000000	0.000000
1	0.00000	0.00000	0.363135	0.437464	0.00000	0.000000	0.000000
2	0.48214	0.48214	0.000000	0.000000	0.00000	0.000000	0.388988
3	0.00000	0.00000	0.000000	0.000000	0.00000	0.614189	0.495524
4	0.00000	0.00000	0.427993	0.000000	0.63907	0.000000	0.000000

  

	like it	movie	not	not good	not like	one
0	0.000000	0.609818	0.000000	0.000000	0.00000	0.00000
1	0.000000	0.437464	0.437464	0.542226	0.00000	0.00000
2	0.000000	0.000000	0.388988	0.000000	0.48214	0.00000
3	0.614189	0.000000	0.000000	0.000000	0.00000	0.00000
4	0.000000	0.000000	0.000000	0.000000	0.00000	0.63907

### EXERCISE-3: Compute Cosine Similarity between 2 Documents

```
In [8]: from sklearn.metrics.pairwise import linear_kernel
```

```
In [9]: doc1 = features[0:1]
doc2 = features[1:2]
score = linear_kernel(doc1, doc2)
print(score)
```

```
[[0.71736783]]
```

```
In [10]: scores = linear_kernel(doc1, features)
print(scores)
```

```
[[1.          0.71736783 0.          0.          0.2166519 ]]
```

```
In [11]: query = "I like this good movie"
qfeature = tfidf.transform([query])
scor = linear_kernel(doc1, features)
print(scor)
```

```
[[1.          0.71736783 0.          0.          0.2166519 ]]
```

## EXERCISE-4: Find Top-N similar documents

**Question-1. Consider the following documents and compute TFIDF values**

```
In [12]: docs=["the house had a tiny little mouse",
               "the cat saw the mouse",
               "the mouse ran away from the house",
               "the cat finally ate the mouse",
               "the end of the mouse story"
            ]
```

**Question-2. Compute cosine similarity between 3rd document ("the mouse ran away from the house") with all other documents. Which is the most similar document?**

```
In [13]: tfidf = TfidfVectorizer(min_df=2, max_df=0.5, ngram_range=(1, 2))
features = tfidf.fit_transform(docs)
print(features)
```

```
(0, 3)      0.7071067811865476
(0, 1)      0.7071067811865476
(1, 2)      0.7071067811865476
(1, 0)      0.7071067811865476
(2, 3)      0.7071067811865476
(2, 1)      0.7071067811865476
(3, 2)      0.7071067811865476
(3, 0)      0.7071067811865476
```

```
In [14]: doc1=features[0:3]
s=linear_kernel(doc1,features)
print(s)
```

```
[[1. 0. 1. 0. 0.]
 [0. 1. 0. 1. 0.]
 [1. 0. 1. 0. 0.]]
```

```
In [15]: scores2 = linear_kernel(doc1, features)
print(scores2)
```

```
[[1. 0. 1. 0. 0.]
 [0. 1. 0. 1. 0.]
 [1. 0. 1. 0. 0.]]
```