

### Exercise 1

```
from sklearn.feature_extraction.text import TfidfVectorizer
import pandas as pd
docs = ("good movie", "not a good movie", "did not like",
        "i like it", "good one")
tfidf = TfidfVectorizer(min_df=2, max_df=0.5,
                        ngram_range=(1, 2))
features = tfidf.fit_transform(docs)
print(features)
df = pd.DataFrame(
    features.toarray(),
    columns = tfidf.get_feature_names())
print(df)
```

### Ex: 2

```
tfidf = TfidfVectorizer(min_df=1, max_df=0.6,
                        ngram_range=(1, 2))
features = tfidf.fit_transform(docs)
print(features)
df = pd.DataFrame(features.toarray())
columns = tfidf.get_feature_names()
print(df)
```

### Ex: 3

```
from sklearn.metrics.pairwise import
    linear_kernel
doc1 = features[0:1]
doc2 = features[1:2]
score = linear_kernel(doc1, doc2)
```

## Natural Language Processing Lab

### Lab3. Computing Document Similarity using VSM

#### EXERCISE-1: Print TFIDF values

```
from sklearn.feature_extraction.text import TfidfVectorizer
import pandas as pd

docs = [
    "good movie", "not a good movie", "did not like",
    "i like it", "good one" ]

# using default tokenizer in TfidfVectorizer
tfidf = TfidfVectorizer(min_df=2, max_df=0.5, ngram_range=(1, 2))
features = tfidf.fit_transform(docs)
print(features)

# Pretty printing
df = pd.DataFrame(
    features.todense(),
    columns=tfidf.get_feature_names())
print(df)
```

#### EXERCISE-2:

1. Change the values of **min\_df** and **ngram\_range** and observe various outputs

#### EXERCISE-3: Compute Cosine Similarity between 2 Documents

```
from sklearn.metrics.pairwise import linear_kernel

# cosine score between 1st and 2nd doc
doc1 = features[0:1]
doc2 = features[1:2]
score = linear_kernel(doc1, doc2)
print(score)

# cosine score between 1st and all other docs
scores = linear_kernel(doc1, features)
print(scores)

# Cosine Similarity for a new doc
query = "I like this good movie"
qfeature = tfidf.transform([query])
scores2 = linear_kernel(doc1, features)
print(scores2)
```

#### EXERCISE-4: Find Top-N similar documents

**Question-1.** Consider the following documents and compute TFIDF values

```
docs=["the house had a tiny little mouse",
      "the cat saw the mouse",
      "the mouse ran away from the house",
      "the cat finally ate the mouse",
      "the end of the mouse story"
      ]
```

**Question-2.** Compute cosine similarity between 3<sup>rd</sup> document ("the mouse ran away from the house") with all other documents. Which is the most similar document?.

**Question-3.** Find Top-2 similar documents for the 3<sup>rd</sup> document based on Cosine similarity values.

```
score = linear_kernel(doc1, featuresdoc2)  
print(scores)
```

```
query = "I like this good movie"  
feature = tfidf.transform([query])  
score = linear_kernel(doc1, features)  
print(score)
```

Ex: 4

Quest: 1

```
docs = ["the house had a tiny little mouse",  
        "the cat saw the mouse",  
        "the mouse ran away from the house",  
        "the cat finally ate the mouse",  
        "the end of the mouse story"]
```

Quest: 2

```
tfidf = TfidfVectorizer(min_df=2, max_df=0.5,  
                        ngram_range=(1, 2))  
features = tfidf.fit_transform(docs)  
print(features)
```

```
doc1 = features[0:3]
```

```
s = linear_kernel(doc1, features)  
print(s)
```

```
scores_2 = linear_kernel(doc1, features)  
print(scores_2)
```